

AI-Assisted Enterprise Database Migration & Modernization

1. Enterprise Problem Statement (Real-World Context)

Enterprises often operate **mission-critical applications on legacy databases** (Oracle, SQL Server, DB2, MySQL) that have evolved over years with:

- Tight coupling between application code and database logic
- Heavy usage of stored procedures, triggers, and vendor-specific SQL
- Poor or outdated documentation
- Large data volumes (TB–PB scale)
- Strict uptime, compliance, and performance SLAs

Typical migration goals include:

- On-prem → Cloud (Oracle → Aurora / PostgreSQL, SQL Server → Azure SQL)
- Commercial → Open-source (Oracle → PostgreSQL)
- Legacy schema redesign for microservices

Current approaches rely heavily on **manual analysis, partial automation, and risky big-bang cutovers**, leading to:

- Long timelines (6–18 months)
- High cost and human error
- Unexpected performance regressions
- Data integrity issues

2. Project Overview

This project aims to build an **AI-agent–driven database migration solution** that automates **analysis, planning, transformation, validation, and explainability** of database migrations using **LangGraph / LangChain**.

The solution will:

- Analyze source DB schema, data, and logic
- Generate a safe, explainable migration strategy
- Automatically transform schemas, queries, and stored logic
- Validate correctness, performance, and data integrity
- Support phased and zero-downtime migration patterns

The PoC will be validated using **open-source databases and schemas** that mimic real enterprise complexity.

3. Objectives

Primary Objectives

- Automate database migration with minimal human intervention
- Preserve **data correctness, performance, and transactional integrity**
- Reduce migration risk and time-to-value
- Provide full auditability and explainability of changes

Secondary Objectives

- Support heterogeneous DB migrations
 - Enable phased / hybrid migration strategies
 - Create reusable AI agents for future migrations
-

4. Scope

In-Scope

- Schema migration (tables, indexes, constraints)
- Data type mapping and transformation
- Stored procedures, functions, triggers migration
- SQL query compatibility refactoring
- Data migration and validation
- Performance benchmarking
- Migration reports and explainability

Out-of-Scope (PoC Phase)

- Application code refactoring (limited to SQL compatibility hints)
 - Production cutover automation
 - Proprietary vendor tools integration (later phase)
-

5. Target Users

- Database Architects
 - Platform & Cloud Migration Teams
 - Data Engineering Teams
 - SRE / DevOps Teams
 - Enterprise Architects
-

6. Supported Migration Scenarios (PoC)

Source	Target
MySQL	PostgreSQL
PostgreSQL	Cloud-managed PostgreSQL
Oracle-like schema	PostgreSQL
SQL Server-like schema	PostgreSQL / Azure SQL

7. High-Level Agent Architecture

LangGraph-based Stateful Agent Flow



8. Functional Requirements

8.1 Source Database Introspection

- Connect to source DB (read-only)
- Extract:

- Schemas, tables, views
 - Indexes, constraints
 - Stored procedures, functions, triggers
 - Row counts, data volumes
 - Build dependency graph across objects
-

8.2 Dependency & Impact Analysis

- Identify:
 - Cross-schema dependencies
 - Circular references
 - Vendor-specific SQL features
 - Classify migration complexity:
 - Low / Medium / High risk objects
 - Detect anti-patterns (e.g., heavy trigger logic)
-

8.3 Migration Strategy Planning

- Decide migration approach:
 - Big-bang
 - Phased (schema → data → logic)
 - Dual-write / sync-based (PoC simulation)
 - Define:
 - Object migration order
 - Rollback checkpoints
 - Validation gates
 - Generate a **human-reviewable migration plan**
-

8.4 Schema Transformation

- Convert:
 - Data types (e.g., NUMBER → NUMERIC)
 - Index types
 - Constraints and defaults
 - Normalize vendor-specific features
 - Apply naming and convention standards
 - Generate target-DB-specific DDL
-

8.5 Stored Logic Conversion

- Translate:
 - Stored procedures

- Functions
 - Triggers
 - Replace unsupported constructs with:
 - Equivalent SQL
 - Application-side logic recommendations
 - Flag manual review items clearly
-

8.6 Data Migration

- Support:
 - Full load
 - Incremental load (PoC simulation)
 - Handle:
 - Large tables
 - Referential integrity
 - Data type transformation
 - Provide resumable migration support
-

8.7 Validation & Reconciliation

- Row count comparison
 - Checksum / hash-based data validation
 - Constraint validation
 - Referential integrity checks
 - Transaction consistency verification
-

8.8 Performance Benchmarking

- Capture baseline metrics:
 - Query execution time
 - Index usage
 - Storage footprint
 - Compare post-migration metrics
 - Flag regressions beyond defined thresholds
-

8.9 Explainability & Reporting

- Generate reports:
 - Objects migrated
 - Transformation rationale
 - Known risks & limitations
- Provide:

- SQL diffs
 - Mapping tables (source → target)
 - Export:
 - PDF / Markdown / JSON formats
-

9. Non-Functional Requirements

9.1 Reliability

- Idempotent migrations
- Safe rollback capability
- Consistent results across runs

9.2 Performance

- PoC scale:
 - 10–50 GB datasets
- No more than ±5–10% performance degradation

9.3 Security

- Read-only source access
- Masking support for sensitive columns
- No credential persistence

9.4 Compliance & Auditability

- Full migration logs
 - Change traceability per DB object
 - Explainable AI decisions
-

10. AI / Agent Requirements

10.1 Framework

- LangGraph for stateful workflows
- LangChain for:
 - SQL understanding
 - Code transformation
 - Tool orchestration

10.2 Agent Capabilities

- SQL AST analysis
- Cross-DB dialect reasoning
- Stateful retries and self-healing

- Tool calling for:
 - DB clients
 - Benchmark execution
 - Validation scripts
-

11. PoC Validation Strategy

11.1 Test Data

- Open-source schemas:
 - Sakila
 - AdventureWorks
 - TPC-H (scaled)
- Introduce:
 - Triggers
 - Stored procedures
 - Vendor-specific SQL constructs

11.2 Success Criteria

- 100% schema migration success
 - $\geq 99.99\%$ data integrity match
 - Acceptable performance variance
 - Clear migration documentation
-

12. Deliverables

- AI-driven migration PoC
 - LangGraph workflow definitions
 - Sample migration reports
 - Demo walkthrough
 - Technical documentation
-

13. Risks & Mitigations

Risk	Mitigation
Complex stored logic	Flag + partial automation
Performance regressions	Automated benchmarking
Data loss	Multi-layer validation

SQL dialect gaps

Agent-driven refactoring

Database Migration / Schema Tools Repositories

1. SQL Server → PostgreSQL Migration Tool

Converts SQL Server schema to PostgreSQL and can generate data migration jobs — great for heterogeneous DB migration PoC.

[https://github.com/dalibo/sqlserver2pgsql GitHub](https://github.com/dalibo/sqlserver2pgsql)

2. Atlas (Schema as Code & Migration Tool)

Declarative schema management and migration tool useful for automated DB schema diff + migration logic.

[https://github.com/ariga/atlas GitHub](https://github.com/ariga/atlas)

3. postgres_migrator (PostgreSQL Migrations CLI)

Generates and applies schema migrations from declarative SQL — useful for migration tooling and workflow tests.

[https://github.com/blainehansen/postgres_migrator GitHub](https://github.com/blainehansen/postgres_migrator)